

Compiling SHACL into SQL

Maxime Jakubowski^{1,2} and Jan Van den Bussche²


¹DBAI @ TU Wien

²DSI @ UHasselt

Shapes Constraint Language (SHACL)

- *Shapes* are constraints on a graph, and consists of two parts:
 - the **shape expression**: describes a constraint on a node
 - the **target declaration**: describes which nodes need to be checked
- Checking whether a shape holds in a graph is called **validation**
- For example:
 - “has an **author**” is a shape expression
 - “every **book**” is a target declaration

```
:BookShape a sh:PropertyShape ;  
  sh:path :author ;  
  sh:minCount 1 .  
  
:BookShape sh:targetClass :Book .
```



Logical SHACL

Shape Expressions

$E ::= p \mid p^- \mid E \cdot E \mid E \cup E \mid E^*$ *(path expressions)*

$\phi ::= \top \mid hasValue(c) \mid eq(E, p) \mid disj(E, p) \mid closed(Q)$ *(shape expressions)*
 $\mid hasShape(s) \mid lessThan(E, p) \mid uniquelang(E)$
 $\mid \geq_n E. \phi \mid \phi \vee \phi \mid \phi \wedge \phi \mid \neg \phi$

- A **shape** is a triple consisting of a **shape name**, **expression**, and **target**
- Target declarations can also be expressed as shapes
- For example: “sh:targetClass :Book” is $\geq_1 \text{ type. } hasValue(:Book)$
- We formulate our previous example in logical SHACL as a triple:
 $(:BookShape, \geq_1 \text{ author}, \geq_1 \text{ type. } hasValue(:Book))$

Validation *is* Querying

- Shape expressions and target declarations are like node-queries:
 - `shape_expr(node)` retrieves all nodes satisfying a shape
 - `target(node)` retrieves all targeted nodes
- Retrieving all ***violating*** nodes, i.e., validation, is the difference query:

$$\text{violation}(\text{node}) = \text{target}(\text{node}) - \text{shape_expr}(\text{node})$$

- SHACL is a powerful language with logical constructs and we can write the violation query as a shape expression itself!
- Our example for `:BookShape` is then:

$$\geq_1 \text{ type.} \textit{hasValue}(:\text{Book}) \wedge \neg \geq_1 \text{ author}$$

Analytical Queries

Hypothesis

SHACL validation is *analytical querying*

- Often non-monotonic queries with aggregation
 - “Qualified Value Shapes” (of the form $\geq_n p. \phi$) are group joins with count aggregation
 - Lessthan constraints are min/max aggregation queries with counting
 - Disjointness constraints and closedness are difference queries
- These are usually formulated in SQL
- Our goal:

Test this hypothesis **out of the box**

Methodology: a translation

- Validation amounts to shape evaluation, so the focus is on the translation of a single shape expression
- Given a shape expression ϕ , we define a SQL query Q_ϕ such that:
For a given graph G , nodes satisfying ϕ in G are exactly the nodes returned by executing Q_ϕ on the relational database representation of G
- We will:
 - Define the relational representation of G
 - Utilize the *negation normal form* of ϕ to obtain an efficient translation
 - Use the SQL database DuckDB "out-of-the-box"

Translation: the database schema

- Tries to stay close to the RDF specification
- Uses “pooling” technique (nodes get an identifier)
- Database schema:
 - **IRIs(Node: int64, Value: string)**
 - **Blanks(Node: int64)**
 - **Literals(Node: int64, Value: string, Type: string, Lang: string)**
 - **Nodes(Node: int64)**
 - **Triples(Subject: int64, Predicate: string, Object: int64)**
 - **Numerics(Node: int64, Value: double)**

Translation: Cardinality Constraints (1)

- (Qualified) min count: $\geq_n p.\phi$

```
SELECT Subject AS Node FROM Triples, (Q( $\phi$ )) AS Q(Node)
WHERE Predicate = p AND Object = Q.Node
GROUP BY Subject
HAVING COUNT(*) >= n
```

- (Qualified) min and max count: $\geq_n p.\phi \wedge \leq_m p.\phi$

```
SELECT Subject AS Node FROM Triples, (Q( $\phi$ )) AS Q(Node)
WHERE Predicate = p AND Object = Q.Node
GROUP BY Subject
HAVING COUNT(*) >= n AND COUNT(*) <= m
```


Translation: Cardinality Constraints (2)

- (Qualified) max count: $\leq_n p.\phi$

SELECT Node **FROM** Nodes
WHERE Node **NOT IN** ($Q_{\geq n+1} p.\phi$)

- (Qualified) universal: $\forall p.\phi$

SELECT Node **FROM** Nodes **WHERE NOT EXISTS** (
 SELECT * **FROM** Triples, (Q_ϕ) **AS** Q(Node)
 WHERE Predicate = p **AND**
 Subject = Node **AND**
 Object **NOT IN** Q)

Translation: Equality Constraint $eq(p, q)$

```
SELECT Node FROM Nodes
WHERE NOT EXISTS
(((      SELECT Object FROM Triples
        WHERE Predicate =  $p$  AND
            Subject = Node
    ) EXCEPT (
        SELECT Object FROM Triples
        WHERE Predicate =  $p$  AND
            Subject = Node
    )) UNION ((
        SELECT Object FROM Triples
        WHERE Predicate =  $q$  AND
            Subject = Node
    ) EXCEPT (
        SELECT Object FROM Triples
        WHERE Predicate =  $p$  AND
            Subject = Node )))
```

Translation: Equality Constraint $\neg eq(p, q)$

```
SELECT Node FROM Nodes
WHERE NOT EXISTS
(((      SELECT Object FROM Triples
        WHERE Predicate = p AND
            Subject = Node
    ) EXCEPT (
        SELECT Object FROM Triples
        WHERE Predicate = p AND
            Subject = Node
    )) UNION ((
        SELECT Object FROM Triples
        WHERE Predicate = q AND
            Subject = Node
    ) EXCEPT (
        SELECT Object FROM Triples
        WHERE Predicate = p AND
            Subject = Node )))
```

```
SELECT Node FROM Nodes
WHERE EXISTS
((      SELECT * FROM Triples
        WHERE Predicate = p AND
            Object NOT IN (
                SELECT Object FROM Triples
                WHERE Subject = Node AND
                    Predicate = q )
    ) UNION (
        SELECT * FROM Triples
        WHERE Predicate = q AND
            Object NOT IN (
                SELECT Object FROM Triples
                WHERE Subject = Node AND
                    Predicate = p )))
```

Translation: everything else

- Our translation covers:
 - All “logical” features:
sh:and, sh:or, sh:not, sh:xone, sh:disjoint, sh:equals, ...
 - Most of the tests on RDF Literals:
sh:lessThan, sh:hasvalue, sh:datatype, sh:pattern, sh:languageIn, ...
- Does not cover:
 - Full path expressions (only inverse)
 - Comparison constraints on values other than numerics
- Full description in the paper, and implemented & available online






Experimental Setup

Three different experiments:

1. Synthetic shapes and data
 - covers many SHACL features
 - simple synthetic data
2. Tyrolian benchmark
 - uses real data
 - mostly conjunctions of datatype tests
3. DBLP data with custom shapes
 - covers more complex shapes
 - uses real data

Synthetic 5mio





Concluding Remarks

- Recent advances in databases are relevant for SHACL validators
- SHACL is like analytical querying
- Our approach runs best if there are many targets
- Expansion to support more complex path expressions
- Expansion to support recursion